

sol1_enhancer_codex

Take-home. Isolation is a process sandbox, not a per-agent tool list.

`bin/role.sh` starts the judge and the doer as their own read-only `codex` `exec` processes. Do not lose the execute bit.

Read `HOW_TO_RUN.md`, `SPEC.md`, and `IMPLEMENTATION_NOTES.md`.

Setup

```
cd solutions/sol1_enhancer_codex  
cp config.json.example config.json  
task clone  
task fence-check
```

`task fence-check` **must not** report `TRUSTED`. If it does, the sandbox is not real. See [IMPLEMENTATION_NOTES.md](#).

Checks with no live poll

```
task test
task checks
python3 .agents/skills/enhancer-loop/scripts/check_fields.py --demo
python3 .agents/skills/enhancer-loop/scripts/check_stop.py --demo
```


One poll. Budget five minutes

```
task create-test-tickets
timeout 420 task run --
task poll-forever --
```

One poll starts three model processes: judge, doer, judge again. A full round that promotes a candidate runs about four minutes.

A run that produces no output for minutes is usually a hang. `IMPLEMENTATION_NOTES.md` lists the two that look identical from outside.

Architecture



Orchestrator writes GitHub. Children are read-only. `task fence-check` must not report TRUSTED.

Testing skill

`.agents/skills/test-sol1-ticket-enhancer/`

Same reset, seed, poll, LGTM ritual. Point SOLUTION at this folder. Read HOW_TO_RUN.md first.

Spillwave Solutions | spillwave.com

Troubleshooting

	SYMPTOM		FIX
TRUSTED		sandbox is not a fence.	IMPLEMENTATION_NOTES.md
	role.sh not found		restore the execute bit
	Hang, no output		timeout 420. Then the notes.
	Closed issue		task reset-test-tickets

Recap

Process sandbox. Children cannot write. Orchestrator alone talks to GitHub.

Expect five minutes, not one.

Spillwave Solutions | spillwave.com